

# Simple Linear Regression

Single Features

1 Feature vs 1 Target: Ex(Study hrs Vs Marks)

# Multiple Liner Regression

Multiple Features

More than 1 Col(Features) vs 1 Target: Ex(Sales, Profit, Loan Amount, GDP)

## Score High(Best), Error(Lowest): Best Model

1: Error's: Loss, Residuals, Cost: Low>>

## Canadas's Per Capita Income(\$)

## Step 1 Load important Modules

```
In [35]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, root_mean_square
import warnings
warnings.filterwarnings('ignore')
print('Done')
```

Done

## Step2: Load Dataset

```
In [2]: file_path = r"C:\Users\Administrator\Downloads\canada_per_capita_income.csv"
df = pd.read_csv(file_path)
print('Done')
```

Done

## EDA:-Exploratory Data Analysis

```
In [3]: df.shape
```

```
Out[3]: (47, 2)
```

```
In [4]: df.sample(5)
```

```
Out[4]:   year  per capita income (US$)
```

19	1989	16426.725480
37	2007	36144.481220
7	1977	7100.126170
34	2004	25719.147150
11	1981	9434.390652

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47 entries, 0 to 46
Data columns (total 2 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   year            47 non-null      int64  
 1   per capita income (US$) 47 non-null      float64 
dtypes: float64(1), int64(1)
memory usage: 884.0 bytes
```

```
In [6]: df.head()
```

```
Out[6]:   year  per capita income (US$)
```

0	1970	3399.299037
1	1971	3768.297935
2	1972	4251.175484
3	1973	4804.463248
4	1974	5576.514583

```
In [7]: df.isna().sum()
```

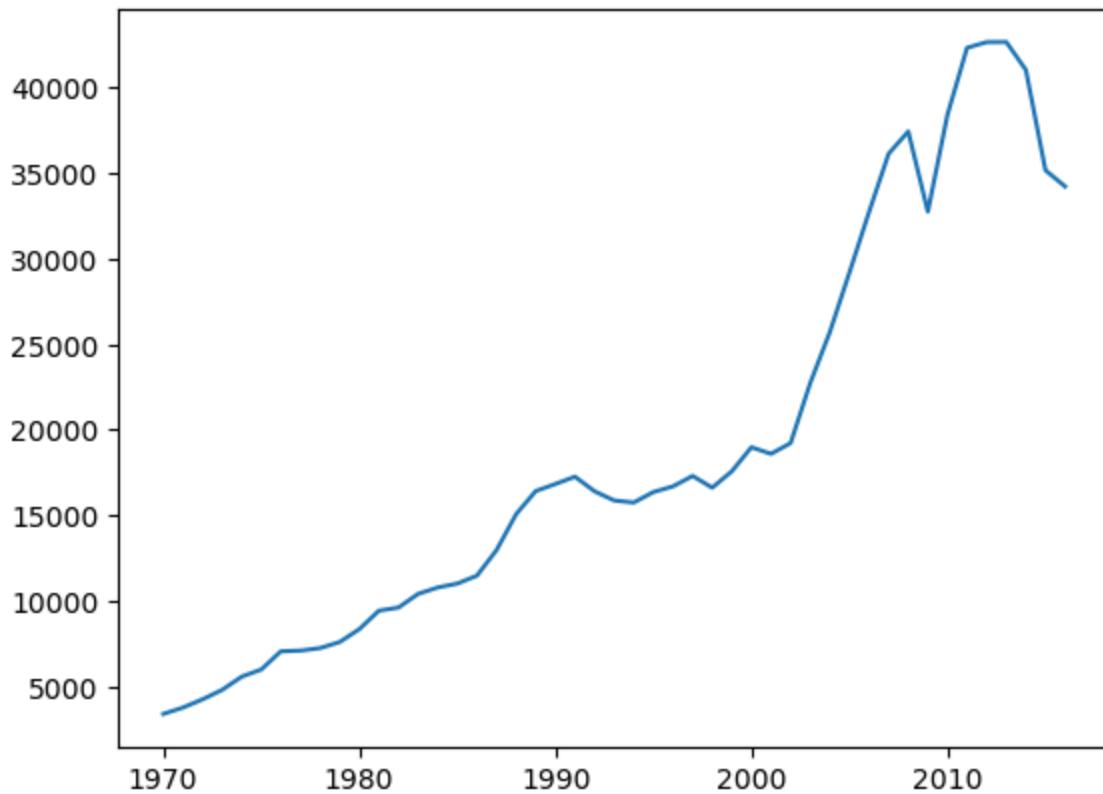
```
Out[7]: year          0
per capita income (US$) 0
dtype: int64
```

```
In [8]: df.describe()
```

Out[8]:

	year	per capita income (US\$)
<b>count</b>	47.000000	47.000000
<b>mean</b>	1993.000000	18920.137063
<b>std</b>	13.711309	12034.679438
<b>min</b>	1970.000000	3399.299037
<b>25%</b>	1981.500000	9526.914515
<b>50%</b>	1993.000000	16426.725480
<b>75%</b>	2004.500000	27458.601420
<b>max</b>	2016.000000	42676.468370

In [9]: `corr = df.corr()`In [10]: `sns.heatmap(corr, annot=True)  
plt.show()`In [11]: `plt.plot(df['year'], df['per capita income (US$)'])  
plt.show()`



## Step3 Dataset Divide X,y

```
In [12]: #3.1  
X = df[['year']]  
y = df['per capita income (US$)']  
  
X.shape, y.shape
```

```
Out[12]: ((47, 1), (47,))
```

```
In [13]: from sklearn.base import RegressorMixin  
#3.2  
  
model = LinearRegression()  
model.fit(X,y)
```

```
Out[13]: ▾ LinearRegression ⓘ ⓘ  
LinearRegression()
```

```
In [14]: y_pred = model.predict(X)
```

```
In [15]: y_pred
```

```
Out[15]: array([-134.55966672,  693.9054085 ,  1522.37048373,  2350.83555895,
 3179.30063417,  4007.7657094 ,  4836.23078462,  5664.69585984,
 6493.16093506,  7321.62601029,  8150.09108551,  8978.55616073,
 9807.02123595,  10635.48631118,  11463.9513864 ,  12292.41646162,
 13120.88153685,  13949.34661207,  14777.81168729,  15606.27676251,
 16434.74183774,  17263.20691296,  18091.67198818,  18920.1370634 ,
 19748.60213863,  20577.06721385,  21405.53228907,  22233.9973643 ,
 23062.46243952,  23890.92751474,  24719.39258996,  25547.85766519,
 26376.32274041,  27204.78781563,  28033.25289085,  28861.71796608,
 29690.1830413 ,  30518.64811652,  31347.11319175,  32175.57826697,
 33004.04334219,  33832.50841741,  34660.97349264,  35489.43856786,
 36317.90364308,  37146.3687183 ,  37974.83379353])
```

```
In [29]: compare_df = df.copy()
```

```
In [32]: compare_df['Y-Pred'] = y_pred
```

```
In [ ]:
```

```
In [33]: compare_df.columns = ['X', 'y', 'y-pred']
```

```
In [34]: compare_df
```

Out[34]:

	<b>x</b>	<b>y</b>	<b>y-pred</b>
<b>0</b>	1970	3399.299037	-134.559667
<b>1</b>	1971	3768.297935	693.905409
<b>2</b>	1972	4251.175484	1522.370484
<b>3</b>	1973	4804.463248	2350.835559
<b>4</b>	1974	5576.514583	3179.300634
<b>5</b>	1975	5998.144346	4007.765709
<b>6</b>	1976	7062.131392	4836.230785
<b>7</b>	1977	7100.126170	5664.695860
<b>8</b>	1978	7247.967035	6493.160935
<b>9</b>	1979	7602.912681	7321.626010
<b>10</b>	1980	8355.968120	8150.091086
<b>11</b>	1981	9434.390652	8978.556161
<b>12</b>	1982	9619.438377	9807.021236
<b>13</b>	1983	10416.536590	10635.486311
<b>14</b>	1984	10790.328720	11463.951386
<b>15</b>	1985	11018.955850	12292.416462
<b>16</b>	1986	11482.891530	13120.881537
<b>17</b>	1987	12974.806620	13949.346612
<b>18</b>	1988	15080.283450	14777.811687
<b>19</b>	1989	16426.725480	15606.276763
<b>20</b>	1990	16838.673200	16434.741838
<b>21</b>	1991	17266.097690	17263.206913
<b>22</b>	1992	16412.083090	18091.671988
<b>23</b>	1993	15875.586730	18920.137063
<b>24</b>	1994	15755.820270	19748.602139
<b>25</b>	1995	16369.317250	20577.067214
<b>26</b>	1996	16699.826680	21405.532289
<b>27</b>	1997	17310.757750	22233.997364
<b>28</b>	1998	16622.671870	23062.462440
<b>29</b>	1999	17581.024140	23890.927515

	x	y	y-pred
30	2000	18987.382410	24719.392590
31	2001	18601.397240	25547.857665
32	2002	19232.175560	26376.322740
33	2003	22739.426280	27204.787816
34	2004	25719.147150	28033.252891
35	2005	29198.055690	28861.717966
36	2006	32738.262900	29690.183041
37	2007	36144.481220	30518.648117
38	2008	37446.486090	31347.113192
39	2009	32755.176820	32175.578267
40	2010	38420.522890	33004.043342
41	2011	42334.711210	33832.508417
42	2012	42665.255970	34660.973493
43	2013	42676.468370	35489.438568
44	2014	41039.893600	36317.903643
45	2015	35175.188980	37146.368718
46	2016	34229.193630	37974.833794

```
In [38]: mae_math = (compare_df['y'] - compare_df['y-pred']).abs().mean()
print('MAE math', mae_math)
```

MAE math 3088.866427771443

```
In [39]: mae = mean_absolute_error(compare_df['y'], compare_df['y-pred'])
print(mae)
```

3088.866427771443

```
In [41]: mse_math = ((compare_df['y'] - compare_df['y-pred'])**2).mean()
print('MSE math', mse_math)
```

MSE math 15462739.061504772

```
In [42]: mse = mean_squared_error(compare_df['y'], compare_df['y-pred'])
print(mse)
```

15462739.061504772

```
In [43]: rmse_math = mse**0.5
print('RMSE math', rmse_math)
```

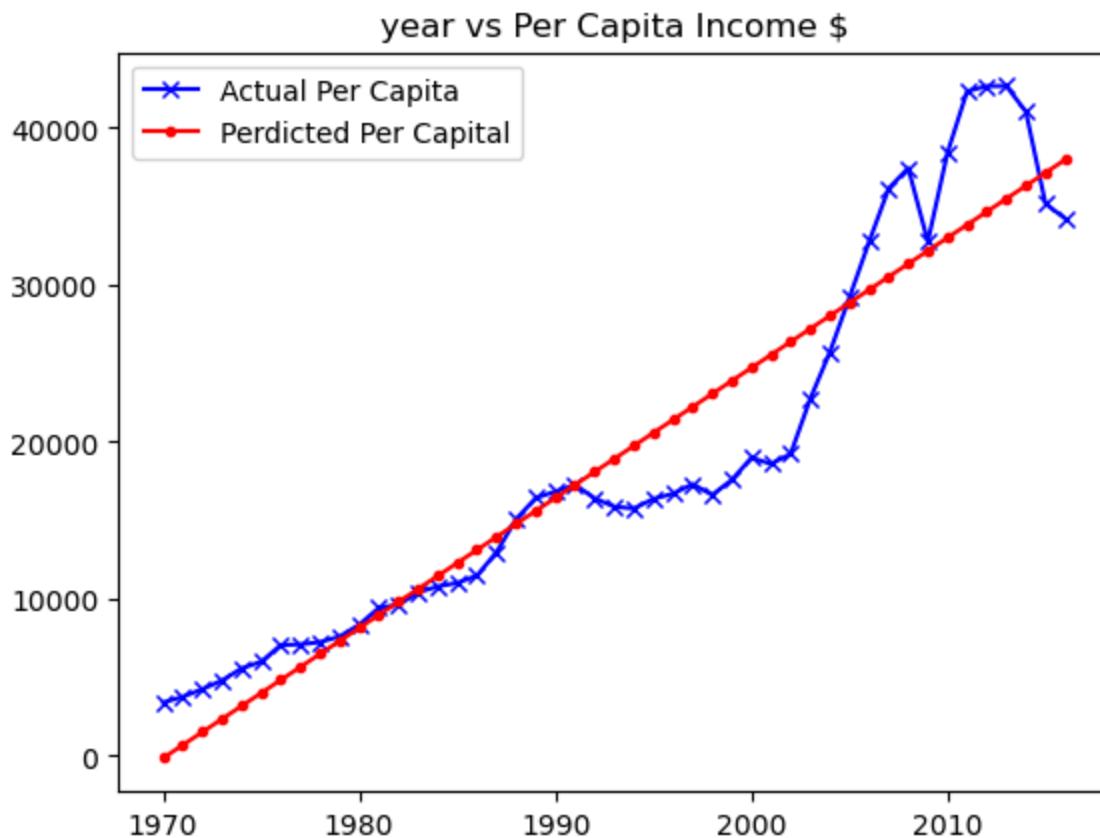
RMSE math 3932.268945723928

```
In [44]: root_mean_squared_error(compare_df['y'], compare_df['y-pred'])
```

```
Out[44]: 3932.268945723928
```

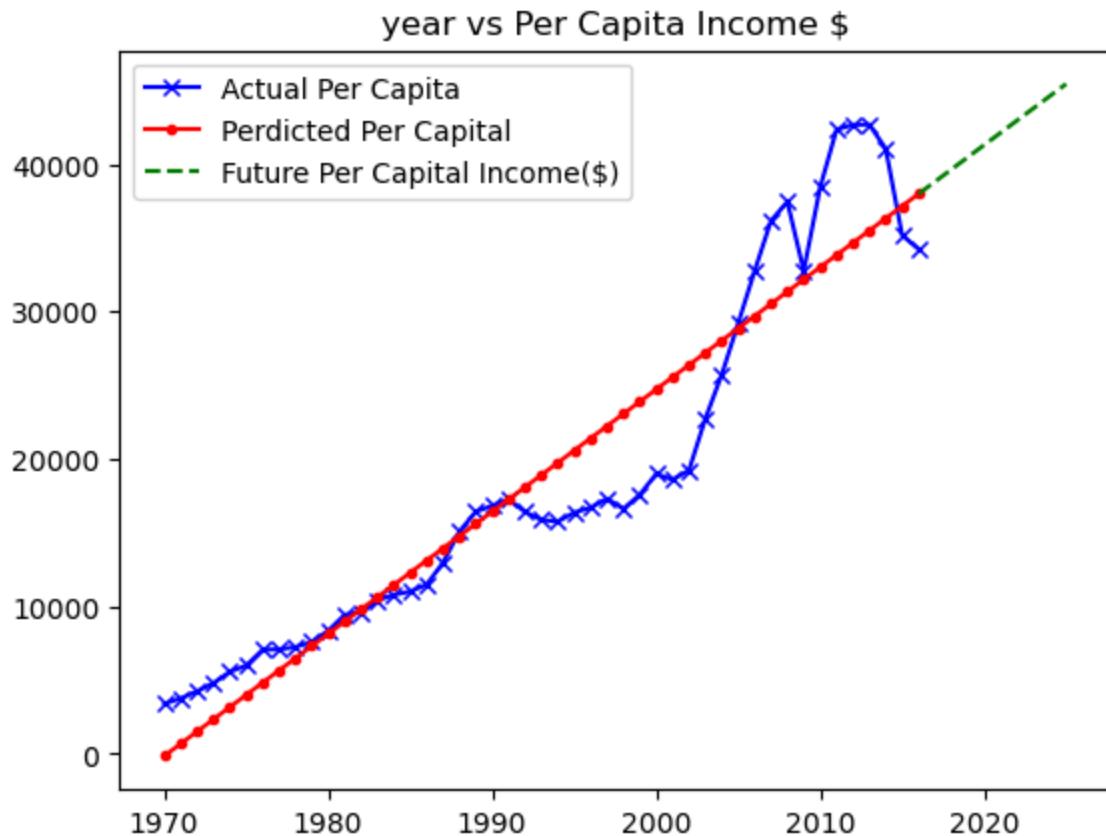
<https://cdn.analyticsvidhya.com/wp-content/uploads/2021/10/42439screenshot-2021-10-26-at-93408-pm-673f1d0b79f01.webp>

```
In [54]: plt.title('year vs Per Capita Income $')
plt.plot(compare_df['X'], compare_df['y'], color = 'b', marker = 'x', label = 'Actual P
plt.plot(compare_df['X'], compare_df['y-pred'], color = 'r', label = 'Predicted Per Ca
plt.legend()
plt.show()
```



```
In [59]: next_years = np.arange(2016,2026).reshape((10,1))
next_years
future_per_capital = model.predict(next_years)
```

```
In [61]: plt.title('year vs Per Capita Income $')
plt.plot(compare_df['X'], compare_df['y'], color = 'b', marker = 'x', label = 'Actual P
plt.plot(compare_df['X'], compare_df['y-pred'], color = 'r', label = 'Predicted Per Ca
plt.plot(next_years,future_per_capital,color = 'g',label ='Future Per Capital Incom
plt.legend()
plt.show()
```



```
In [62]: future_per_capital
```

```
Out[62]: array([37974.83379353, 38803.29886875, 39631.76394397, 40460.22901919,
   41288.69409442, 42117.15916964, 42945.62424486, 43774.08932009,
   44602.55439531, 45431.01947053])
```

```
In [ ]:
```